Real-Time Detection of Alcohol Impairment in Drivers Using Camera-Based and Minimal Sensor ADAS Integration

1. Executive Summary

The global automotive industry stands at a critical juncture where the integration of Artificial Intelligence (AI) into vehicle safety systems is transitioning from passive monitoring to active intervention. While Advanced Driver Assistance Systems (ADAS) have achieved high reliability in detecting driver drowsiness through established metrics like PERCLOS (Percentage of Eye Closure), a significant technological and safety gap remains in the non-intrusive detection of alcohol impairment. Intoxicated driving accounts for a substantial proportion of traffic fatalities worldwide—approximately 30% in the United States alone—yet current vehicular countermeasures typically rely on intrusive breathalyzers or post-incident forensic analysis. This research report presents a comprehensive, exhaustive technical design for a Real-Time Intoxicated Driver Detection System (IDDS) that functions within the constraints of modern automotive embedded hardware.

The proposed system addresses the challenge of accurately predicting intoxication without biochemical sensors by leveraging a "Local-Edge" sensor fusion architecture. This architecture synthesizes data from two primary, non-intrusive streams: camera-based computer vision and vehicle telemetry. The vision module moves beyond simple fatigue detection to analyze complex oculomotor signatures indicative of central nervous system depressants, specifically Horizontal Gaze Nystagmus (HGN), Gaze Transition Entropy (GTE), and the physiological alcohol flush reaction detected via remote photoplethysmography (rPPG). Concurrently, the vehicle behavior module analyzes longitudinal and lateral control dynamics, utilizing Approximate Entropy (ApEn) to quantify the specific "disorder" in steering signals that characterizes cognitive disinhibition, distinguishing it from the "absence" of steering control associated with fatigue.

A core innovation detailed in this report is the application of Dempster-Shafer Theory (DST) for sensor fusion. Unlike Bayesian approaches, which struggle to model "ignorance" or sensor unavailability (e.g., camera blindness due to glare), DST allows the system to explicitly manage uncertainty and conflicting evidence between the vision and telemetry streams. The fusion engine integrates these inputs into a Hybrid AI Model, combining Deep Learning (CNNs and Transformers) for high-dimensional feature extraction with Explainable Machine Learning (Random Forests) for transparent, safety-critical decision-making.

This document serves as a blueprint for the implementation of this system, detailing the mathematical foundations of the entropy-based metrics, the specific architectures of the neural networks, the design of synthetic training environments using the CARLA simulator, and the hardware specifications required for deployment on edge platforms such as the NVIDIA Jetson Orin Nano or Texas Instruments TDA4VM. The analysis confirms that a multi-modal, camera-and-telemetry fusion approach can achieve high-recall intoxication detection, offering a viable path for OEMs to meet emerging regulatory safety mandates (such as the EU General Safety Regulation) while preserving driver privacy through local processing.[1]

2. Introduction: The Imperative for Advanced Impairment Monitoring

2.1 The Epidemiology of Impairment and the Technological Gap

Traffic accidents resulting from driver impairment remain a persistent public health crisis. Despite decades of public awareness campaigns and strict legal penalties, alcohol-impaired driving continues to claim thousands of lives annually. The National Highway Traffic Safety Administration (NHTSA) data indicates that widely deployed safety systems have plateaued in their effectiveness against drunk driving because they are largely reactive—mitigating crash severity rather than preventing the risky behavior itself.[3]

Current ADAS technologies, such as Lane Keeping Assist (LKA) and Automatic Emergency Braking (AEB), operate on the assumption of a rational driver who may occasionally err. They correct the vehicle's trajectory but do not assess the driver's fitness to operate the machine. This limitation becomes acute in the context of SAE Level 2 and Level 3 automation, where the "handover" of control from the vehicle to the driver requires the driver to be cognitively capable of regaining situational awareness immediately. An intoxicated driver, suffering from delayed reaction times and cognitive tunneling, represents a catastrophic failure mode for these semi-autonomous systems.

Existing countermeasures for alcohol detection are predominantly intrusive. Alcohol Interlock Systems (breathalyzers connected to the ignition) are highly effective but are typically mandated only for repeat offenders due to their invasiveness and operational friction. There is no standard, mass-market solution that passively monitors for intoxication with the same ubiquity and acceptance as drowsiness detection systems. This report addresses this gap by proposing a system that utilizes sensors already present in modern "smart" vehicles—specifically the Driver Monitoring System (DMS) camera and the Controller Area Network (CAN) bus—to infer intoxication through behavioral proxies.

2.2 The Challenge of Differentiation: Drunk vs. Drowsy vs. Distracted

A central technical challenge in designing passive impairment detection is distinguishing between different types of driver states. Drowsiness, distraction, and intoxication share overlapping symptoms, such as lane deviation and delayed reaction times, but their underlying neurological mechanisms and behavioral signatures are distinct. A system that conflates these states risks applying the wrong intervention strategy (e.g., a loud wake-up alarm might be effective for a drowsy driver but could induce a panic steering maneuver in a disinhibited, intoxicated driver).

- **Drowsiness:** Characterized by a gradual withdrawal of attention and motor input. The driver fights to stay awake, leading to long periods of inactivity (microsleeps) followed by jerky corrections. Ocularly, it is defined by slow eyelid closure (high PERCLOS) and reduced blink frequency.[5]
- **Cognitive Distraction:** Characterized by "looking but not seeing." The driver may be visually fixated on the road but cognitively disengaged (e.g., daydreaming or talking on a hands-free device). This often results in "visual tunneling," where the gaze distribution narrows significantly.[7]
- **Alcohol Intoxication:** Characterized by cognitive disinhibition and sensorimotor degradation. The driver often believes they are in control but exhibits poor fine motor skills. This results in "sawtooth" steering profiles (over-correction), reduced peripheral scanning due to alcohol myopia, and specific involuntary eye movements like nystagmus that are not present in drowsiness or distraction.[9]

The objective of the proposed system is to disentangle these states by analyzing the *quality* and *entropy* of the control signals, rather than just the magnitude of the errors.

2.3 The "Local-Edge" Computing Paradigm

The architectural philosophy of this research is grounded in "Local-Edge" computing. Early telematics solutions relied on cloud-based analysis, where data snippets were uploaded to central servers for processing. However, cloud dependence is structurally unsound for primary safety systems due to latency, reliability, and privacy constraints.

- **Latency:** A round-trip delay of 500ms—common in congested cellular networks—translates to a vehicle traveling 14 meters at highway speeds. This delay is unacceptable for real-time intervention.[1]
- **Reliability:** Safety systems must function in tunnels, rural areas, and cellular dead zones. A dependency on 4G/5G connectivity introduces a point of failure that compromises the system's Automotive Safety Integrity Level (ASIL) rating.
- **Privacy:** The transmission of video data from inside the cabin raises profound privacy concerns and triggers strict regulatory compliance requirements under GDPR, CCPA, and other data protection laws.

Therefore, the proposed IDDS is designed to run entirely on the vehicle's embedded compute platform. This requires highly optimized algorithms capable of running on SoCs like the NVIDIA Jetson Orin Nano or TI TDA4VM, balancing high-performance inference with strict thermal and power envelopes.[1]

3. Physiological and Oculomotor Detection Modalities

The vision module serves as the primary sensor for detecting the physiological precursors of intoxication. Alcohol acts as a central nervous system depressant, affecting the vestibular system and oculomotor control in ways that are distinct from fatigue. By leveraging high-frame-rate Global Shutter cameras (operating in the Near-Infrared spectrum), the system can extract these subtle biomarkers.

3.1 Horizontal Gaze Nystagmus (HGN)

Horizontal Gaze Nystagmus is the involuntary jerking of the eyeball as it gazes to the side. It is the most reliable field sobriety test used by law enforcement, correlating highly with Blood Alcohol Concentration (BAC) levels above 0.08%.[9]

3.1.1 Neurological Mechanism

The smooth pursuit of a moving target requires the brain to integrate signals from the vestibular system (inner ear) and the visual cortex to maintain foveal fixation. Alcohol impairs the neural integrator in the brainstem, which is responsible for holding the eye in an eccentric position against the elastic restoring forces of the eye muscles. When intoxicated, the eye drifts back toward the center (centripetal drift) and the brain compensates with a corrective saccade (jerk) to snap the gaze back to the target. This cycle repeats rapidly, creating the nystagmus effect.[5]

3.1.2 Passive Detection Algorithm

Traditionally, HGN testing involves a police officer moving a stimulus (pen or light) horizontally. In an automated vehicle context, the system must detect HGN *passively* by exploiting natural driving behaviors, such as checking side mirrors or tracking passing vehicles.[13]

The detection algorithm focuses on three key clues defined in the Standardized Field Sobriety Test (SFST):

1. **Lack of Smooth Pursuit:** As the eye tracks a moving object (or as the head turns while maintaining gaze on a fixed point), the velocity profile of the pupil is analyzed. A sober eye moves smoothly; an intoxicated eye moves in a series of micro-saccades. The algorithm calculates the variance of the pupil's velocity vector v_{pupil}. If the variance exceeds a threshold T_{jerk} during a pursuit maneuver, it logs a "Lack of Smooth Pursuit" event.[5]
2. **Distinct and Sustained Nystagmus at Maximum Deviation:** When the driver looks at the side mirror (typically $> 45^\circ$ eccentricity), the eye is held at maximum deviation. The system monitors the pupil position for high-frequency (2-4 Hz) oscillation. A Fast Fourier Transform (FFT) is applied to the horizontal pupil position signal $x(t)$ over a short

window. A spectral peak in the nystagmus frequency band indicates impairment.[9]

3. **Onset Prior to 45 Degrees:** The angle of gaze θ_{gaze} at which the jerking begins correlates with BAC. The formula typically used is $BAC \approx 50 - \theta_{onset}$. If nystagmus is detected before the eye reaches 45 degrees of deviation, it suggests a high BAC.[9]

3.2 Gaze Entropy and Visual Scanning Behavior

While HGN focuses on the *mechanics* of the eye, gaze entropy analyzes the *strategy* of visual attention. Intoxicated drivers suffer from "alcohol myopia," a narrowing of the attentional field, leading to staring or, conversely, erratic scanning due to disinhibition.

3.2.1 Stationary Gaze Entropy (SGE)

SGE measures the spatial dispersion of the driver's gaze across the visual field. The driver's field of view is divided into Regions of Interest (ROIs): Road Center, Dashboard, Left Mirror, Right Mirror, Infotainment. SGE is calculated using Shannon's entropy formula:

$$H_S = -\sum_{i=1}^{K} p_i \log_2 p_i$$

Where p_i is the probability (proportion of time) the gaze dwells in ROI i.

- **Interpretation:** A sober driver actively scans the environment (checking mirrors, speed, road), resulting in a consistently high SGE. An intoxicated driver typically exhibits **lower SGE**, indicating visual tunneling (fixating on the road center to maintain control) or a "fixed stare".[10] However, extremely disinhibited drivers might show random, non-strategic scanning, which must be differentiated from active scanning via transition analysis.

3.2.2 Gaze Transition Entropy (GTE)

GTE measures the unpredictability of the *sequence* of gaze fixations. It is calculated using the conditional entropy of the Markov chain formed by the gaze transitions between ROIs:

$$H_T = -\sum_{i=1}^{K} p_i \sum_{j=1}^{K} p(j|i) \log_2 p(j|i)$$

Where $p(j|i)$ is the probability of transitioning to ROI j given the current fixation is in ROI i.

- **Interpretation:** Sober drivers have structured, predictable scanning patterns (e.g., Road -> Mirror -> Road), resulting in lower GTE. Intoxicated drivers exhibit **higher GTE** (more random, disordered transitions) or significantly disrupted patterns due to cognitive impairment affecting top-down attention control.[17] The combination of Low SGE (tunneling) and High GTE (disordered scanning when they do move) is a strong signature of intoxication.[10]

3.3 Remote Photoplethysmography (rPPG) and Facial Flushing

Alcohol induces vasodilation, causing the capillaries near the skin surface to expand. This leads to the "alcohol flush reaction," particularly in individuals with ALDH2 deficiency, but is measurable as increased blood perfusion in the general population.

3.3.1 Principle of rPPG

rPPG extracts the cardiac pulse signal from video by measuring the subtle variations in light absorption by hemoglobin in the skin. Green light (approx. 530nm) is absorbed most strongly by hemoglobin.

- **Algorithm:**
  1. **ROI Selection:** The system identifies stable skin regions (forehead, cheeks) using facial landmarks, avoiding eyes and mouth to minimize motion noise.[19]
  2. **Signal Extraction:** The spatial average of the green channel intensity $G(t)$ is calculated for each frame.
  3. **Signal Separation:** Blind Source Separation (BSS) techniques like Independent Component Analysis (ICA) or advanced algorithms like **Plane-Orthogonal-to-Skin (POS)** are used to separate the pulse signal from noise caused by head movements and lighting changes. POS projects the RGB signals onto a plane orthogonal to the skin tone vector to isolate the pulsatile component.[21]

3.3.2 Alcohol Biomarkers from rPPG

- **Heart Rate (HR):** Alcohol consumption typically elevates the resting heart rate.
- **Heart Rate Variability (HRV):** Alcohol suppresses the parasympathetic nervous system, leading to a reduction in HRV. The intervals between heartbeats become more regular (metronomic) compared to the healthy variability of a sober individual.[23]
- **DC Component Analysis:** The "DC" (steady) component of the rPPG signal corresponds to the overall blood volume. A significant increase in the DC component of the green channel relative to the baseline suggests vasodilation (facial flushing).[25]

4. Vehicle Dynamics and Telemetry Modalities

While vision captures the physiological state, vehicle telemetry captures the functional manifestation of impairment in driving performance. Using the CAN bus, the system accesses high-frequency (50Hz+) data on steering, speed, and pedal actuation.

4.1 Steering Entropy (ApEn): The Measure of Control Chaos

Steering Entropy is widely regarded as the most sensitive metric for detecting driver impairment because steering is a continuous control task.

4.1.1 Approximate Entropy (ApEn) Calculation

ApEn quantifies the irregularity of a time series. For steering, it measures how much the driver's steering inputs deviate from a "smooth" or predicted path.

- **Step 1: Prediction Model.** A second-order Taylor series expansion or an Autoregressive (AR) model is used to predict the steering angle $\theta_{pred}(t+1)$ based on the recent history of steering inputs. This represents the "ideal" smooth steering required for the road geometry.

$$\theta_{pred}(n+1) = \theta(n) + \theta'(n)\Delta t + 0.5\theta''(n)\Delta t^2$$

- **Step 2: Residual Calculation.** The prediction error (residual) is the difference between the actual steering angle and the predicted angle: $r(n) = \theta_{actual}(n) - \theta_{pred}(n)$.
- **Step 3: Entropy Computation.** ApEn is calculated on the time series of residuals $r(n)$.

$$ApEn(m, r, N) = \Phi^m(r) - \Phi^{m+1}(r)$$

  Where m is the pattern length (e.g., 2), r is the tolerance threshold (e.g., 0.2 * SD), and Φ is the frequency of patterns.

4.1.2 Differentiating States via Entropy

- **Sober/Alert:** The driver makes frequent, small, smooth micro-corrections. The residuals are small and random, but the overall control is tight. ApEn is moderately high but bounded.[27]
- **Drowsy:** The driver stops steering (micro-sleep), leading to near-zero entropy (highly predictable/regular), followed by a sudden, large correction (spike). The overall ApEn drops significantly (< 0.4).[1]
- **Intoxicated:** The driver's motor control is degraded. They exhibit "bang-bang" control—overcorrecting, waiting too long, then overcorrecting back. The residuals are large and erratic. The ApEn is often **significantly higher** or exhibits a distinct "jagged" distribution compared to the smooth randomness of a sober driver.[27]

4.2 Steering Wheel Reversal Rate (SWRR)

SWRR counts the frequency of steering reversals (changes in direction) that exceed a certain gap angle (e.g., 3 degrees).

- **Intoxication Signature:** Drunk drivers often exhibit a **high SWRR**. Due to compromised feedback loops, they constantly oscillate around the lane center, correcting left, then right, in a continuous, high-amplitude manner.[29]
- **Comparison:** Drowsy drivers typically show a **low SWRR** (drifting without correcting) until a critical lane departure event occurs.[31]

4.3 Lateral and Longitudinal Control

- **Standard Deviation of Lateral Position (SDLP):** This measures the weaving of the car. Alcohol causes a continuous, sinusoidal weaving pattern, significantly increasing SDLP even at low BACs (0.05%).[32] Drowsiness tends to cause gradual drifts rather than weaving.
- **Pedal Dynamics:** Intoxication impairs the fine motor control of the feet. This manifests as **Throttle Oscillation** (inability to hold a steady speed) and **Brake Jitter** (erratic pressure application during stops). Frequency domain analysis of the throttle position sensor can reveal these high-frequency tremors.[1]

5. Sensor Fusion and Artificial Intelligence Architecture

To achieve high robustness, the system must fuse the vision and telemetry streams. A simple "voting" mechanism is insufficient because sensors often disagree (e.g., sunglasses block the camera, or a straight road provides no steering data). We employ Dempster-Shafer Theory (DST) for fusion and a Hybrid AI architecture for classification.

5.1 Dempster-Shafer Theory (DST) for Uncertainty Management

DST is a generalization of Bayesian theory that allows for the assignment of belief to "sets" of hypotheses, including the set of "Uncertainty" or "Ignorance."

5.1.1 The Frame of Discernment

The system defines a set of mutually exclusive states:
$\Theta = \{Sober(S), Drowsy(Dr), Drunk(Dk)\}$.

5.1.2 Mass Function Assignment

Each sensor module acts as an "expert" assigning a basic belief mass (m) to subsets of Θ.

- **Vision Module (m_v):**
  - If HGN is detected: $m_v(\{Dk\}) = 0.7$, $m_v(\Theta) = 0.3$.
  - If PERCLOS > 0.2: $m_v(\{Dr\}) = 0.8$, $m_v(\Theta) = 0.2$.

- If Face not detected (occlusion): $m_v(\Theta) = 1.0$ (Total ignorance).

- **Telemetry Module (m_t):**
  - If ApEn is High & SWRR is High: $m_t(\{Dk\}) = 0.6$, $m_t(\{Dr, Dk\}) = 0.2$ (Could be erratic drowsiness), $m_t(\Theta) = 0.2$.
  - If ApEn is Low: $m_t(\{Dr\}) = 0.7$, $m_t(\Theta) = 0.3$.

5.1.3 Dempster's Rule of Combination

The core fusion logic combines the masses from Vision and Telemetry to calculate a joint belief:

$$m_{1,2}(A) = \frac{\sum_{B \cap C = A} m_1(B) m_2(C)}{1 - K}$$

Where K is the conflict coefficient: $K = \sum_{B \cap C = \emptyset} m_1(B) m_2(C)$.

- **Handling Conflict:** If Vision says "Drunk" (HGN) and Telemetry says "Sober" (Smooth driving), K increases. DST normalizes the result by $1 - K$, effectively highlighting the agreement where it exists or identifying a high-conflict state that requires the system to hold its decision rather than false-alarming. This prevents the system from defaulting to "Sober" just because one sensor is noisy.[1]

5.2 Hybrid AI Model Architecture

The fusion engine feeds into a two-stage learning architecture designed for compliance with ISO 26262 (Functional Safety), which discourages "black box" end-to-end models for safety-critical decisions.

5.2.1 Stage 1: Deep Learning for Perception (Feature Extraction)

- **Vision Stream (CNN):** A lightweight Convolutional Neural Network (CNN), such as **MobileNetV3-Small**, is used for facial feature extraction. It is pre-trained on datasets like FaceHQ and fine-tuned on the DIF (Dataset of Intoxicated Faces) to recognize visual cues like eyelid droop and facial flushing. It outputs a feature vector (embeddings) rather than a classification.[1]
- **Telemetry Stream (LSTM/Transformer):** A **Long Short-Term Memory (LSTM)** network or a small **Transformer** block processes the time-series CAN data. This captures temporal dependencies—for example, recognizing the sequence of "drift -> jerk -> overcorrect" which is characteristic of impairment. It outputs a temporal embedding vector.[37]

5.2.2 Stage 2: Explainable Decision Trees

- **Random Forest / XGBoost:** The high-level features (e.g., HGN Score from Vision, ApEn Score from Telemetry, rPPG HRV) are fed into a Random Forest classifier.
- **Why Trees?** Tree-based models provide **feature importance (SHAP values)**. If the system triggers an alarm, it can generate a trace log: *"Alert Triggered. Reason: HGN Index > 0.8 AND Steering Entropy > Threshold."* This explainability is crucial for legal validation and debugging, unlike an end-to-end Neural Network which outputs a decision without rationale.[1]

5.3 Classification Decision Matrix

The fusion logic effectively implements a multi-dimensional decision boundary:

Feature Metric	**Sober State**	**Drowsy State**	**Drunk State**
PERCLOS	Low (< 0.15)	**High (> 0.20)**	Low / Normal
HGN (Nystagmus)	Absent	Absent	**Present (High Conf)**
Gaze Entropy (SGE)	High (Active)	Low (Staring)	**Low (Tunneling)**
Steering Entropy	Medium (Controlled)	**Low (Passive)**	**High (Chaotic)**
Steering Reversals	Moderate	Low	**High (Jerky)**
HRV (rPPG)	High (Variable)	Moderate	**Low (Metronomic)**

Table 1: Feature Matrix for Driver State Classification [1]

6. Simulation, Datasets, and Validation Strategy

Developing this system requires a rigorous data strategy. Since collecting real-world drunk driving data is dangerous and legally restricted, a hybrid approach using high-fidelity simulation and specialized datasets is required.

6.1 Synthetic Data Generation: The CARLA Simulator

The CARLA simulator (an open-source simulator for autonomous driving research) is used to generate the "Drunk" telemetry dataset.

- **Simulating Impairment:** Python scripts utilizing the CARLA Python API inject specific noise models into the ego vehicle's control loop to mimic intoxication.
  - **Delayed Reaction:** A buffer queue adds a 200ms-500ms delay to steering and brake inputs, simulating cognitive slowing.[39]
  - **Motor Noise:** Gaussian noise is added to the steering angle to simulate poor fine motor control (tremors).
  - **Visual Tunneling:** A post-processing shader blurs the peripheral regions of the driver's camera view, forcing the simulated agent (or human-in-the-loop) to ignore peripheral hazards, generating realistic "drunk" gaze data.[41]
- **Scenario Design:** The simulation runs scenarios specifically designed to elicit impairment markers, such as long highway curves (testing lateral control) and sudden obstacle appearances (testing reaction time and braking jitter).[43]

6.2 Real-World Datasets

- **DIF (Dataset of Intoxicated Faces):** Contains audio-visual clips of intoxicated individuals. Used to train the CNN to recognize facial flushing and heavy eyelids.[44]
- **Drbife / ADDS:** Recent multimodal datasets that include physiological and sensor data for drowsiness and intoxication.[45]
- **NTHU-DDD:** A definitive dataset for drowsiness detection (yawning, nodding) used to train the "Drowsy" class to ensure the model can distinguish it from "Drunk".[1]

6.3 Validation Metrics

The system is evaluated based on its ability to minimize False Negatives (missing a drunk driver).

- **Recall (Sensitivity):** The primary metric. Target > 90% for the "Drunk" class.
- **Precision:** Target > 80% to prevent "alarm fatigue" (annoying the driver with false alarms).
- **Latency:** The system must converge on a decision within a rolling window of 60 seconds of driving data.[1]

7. Hardware Implementation and Embedded Design

The transition from research code to a vehicle-deployable system requires selecting hardware that meets automotive standards for thermal management, vibration resistance, and power efficiency.

7.1 System-on-Chip (SoC) Selection

Based on the computational load of running simultaneous CNN (Vision) and LSTM (Telemetry) models, three hardware tiers are evaluated:

1. **NVIDIA Jetson Orin Nano (Recommended for Prototype):**
   - **Specs:** 6-core ARM CPU, Ampere GPU (1024 CUDA cores), 40 TOPS AI performance.
   - **Advantage:** Supports the full NVIDIA DeepStream SDK for optimized video pipelines. TensorRT can quantize models to INT8 precision, drastically reducing latency.
   - **Power:** Configurable 7W-15W mode fits within the auxiliary power budget of most vehicles.[1]
2. **Texas Instruments TDA4VM (Production/Safety):**
   - **Specs:** Dual C7x DSPs + Matrix Multiply Accelerator (MMA), 8 TOPS.
   - **Advantage:** ASIL-D certified for functional safety. The DSP architecture is vastly more efficient per watt for vision tasks (optical flow, rPPG) than a general-purpose GPU. Ideal for OEM integration.[1]
3. **NXP i.MX 8M Plus (Mass Market):**
   - **Specs:** Integrated NPU (2.3 TOPS).
   - **Advantage:** Ultra-low power (2-4W), enabling fanless, sealed enclosures. Suitable for cost-sensitive implementations where slightly lower frame rates (30fps) are acceptable.[1]

7.2 Sensor Specifications

- **Camera:** A **Global Shutter** sensor (e.g., OmniVision OV2311 or Onsemi AR0234) is non-negotiable. Rolling shutter sensors introduce "jello effects" due to vehicle vibration, which corrupts the precise geometric calculations required for HGN and Gaze tracking.
- **Illumination:** A synchronized **940nm IR LED** ring. 940nm is invisible to the human eye, preventing driver distraction at night, whereas 850nm emits a faint red glow.[1]
- **IMU:** An automotive-grade 6-axis IMU (e.g., ISM330DHCX) is integrated to provide ground-truth vehicle motion, allowing the system to filter out road bumps from the steering entropy analysis.[1]

8. Ethical, Legal, and Future Considerations

8.1 Privacy and Data Security

To ensure public acceptance and regulatory compliance (GDPR), the system must adhere to **Privacy by Design**.

- **Edge Processing:** Video data is processed in Random Access Memory (RAM) and discarded immediately. No video is ever recorded or transmitted to the cloud. Only the metadata (e.g., "Impairment Level: High") is logged.
- **Data Security:** The system utilizes a Trusted Platform Module (TPM) to store cryptographic keys. Any alert signal sent to the vehicle's CAN bus (to trigger a warning)

or to a fleet management system is digitally signed, preventing "spoofing" attacks.[1]

8.2 Functional Safety (ISO 26262)

The IDDS is a safety-critical monitoring system. It must be designed to avoid "Safety Hazards" such as false positives that disable the vehicle at high speed.

- **Safety Shell:** The AI model is wrapped in a deterministic "Safety Shell" (coded in MISRA C). This logic gate validates the AI's output against vehicle speed and location. For example, if the AI predicts "Drunk" with low confidence, the Safety Shell ensures the response is a "Soft Warning" (chime) rather than a "Hard Intervention" (limiting speed), preserving the driver's ultimate authority.[1]

8.3 Future Outlook: V2X and Personalization

- **V2X Integration:** In the future, a detected impaired driver could trigger a Vehicle-to-Everything (V2X) broadcast message—"Erratic Driver detected at GPS X,Y"—alerting nearby autonomous vehicles to widen their safety margins proactively.
- **Personalization:** Using Federated Learning, the system could learn the specific "sober baseline" steering and gaze patterns of the vehicle's owner over time. This would make the detection of deviations (entropy changes) significantly more sensitive and accurate for that specific driver.[1]

9. Conclusion

This research report demonstrates that a robust, real-time Intoxicated Driver Detection System can be constructed without relying on intrusive biochemical sensors. By synthesizing the physiological precision of computer vision (HGN, Gaze Entropy, rPPG) with the robust dynamics of vehicle telemetry (Steering Entropy, SWRR), the system effectively fills the safety gap between drowsiness detection and alcohol impairment monitoring. The proposed "Local-Edge" architecture, powered by Dempster-Shafer sensor fusion and hybrid AI modeling, offers a scalable, privacy-centric solution that meets the rigorous demands of the automotive industry. As regulatory pressure mounts for "impairment detection" beyond simple fatigue, this multi-modal fusion approach represents the state-of-the-art in passive vehicle safety.

Table 2: Proposed System Bill of Materials (BOM) & Specifications

Component	Specification	Purpose	Source

Compute	NVIDIA Jetson Orin Nano (8GB)	AI Inference (Vision + Fusion)	[1]
Camera	Onsemi AR0234 (Global Shutter)	HGN & Gaze Tracking (60fps)	[1]
Lens/Filter	M12 Lens + 940nm Bandpass	Night Vision / Glare Removal	[49]
IMU	STMicro ISM330DHCX (6-axis)	Vehicle Dynamics / Vibration Filter	[1]
Connectivity	Quectel RM520N (5G)	Fleet Alerts / V2X Broadcast	[1]
Security	Infineon OPTIGA TPM	Secure Boot / Data Signing	[1]

Citations

[1]

Works cited

1. On-Vehicle_Local_AI_Driver_State_Detection_Report.pdf
2. An Intelligent Online Drunk Driving Detection System Based on Multi-Sensor Fusion Technology - MDPI, accessed on February 7, 2026, https://www.mdpi.com/1424-8220/22/21/8460
3. Assessing the Feasibility of Vehicle-based Sensors to Detect Drowsy Driving - NHTSA, accessed on February 7, 2026, https://www.nhtsa.gov/sites/nhtsa.dot.gov/files/811886-assess_veh-based_sensors_4_drowsy-driving_detection.pdf
4. USING VEHICLE-BASED SENSORS OF DRIVER BEHAVIOR TO DETECT ALCOHOL IMPAIRMENT - Research, accessed on February 7, 2026, https://www-nrd.nhtsa.dot.gov/departments/esv/22nd/files/22ESV-000322.pdf
5. Assessment of Driver Monitoring Systems for Alcohol Impairment Detection and Level 2 Automation - ROSA P, accessed on February 7, 2026, https://rosap.ntl.bts.gov/view/dot/77647/dot_77647_DS1.pdf
6. Classification of Fatigued and Drunk Driving Based on Decision Tree Methods: A Simulator Study - PMC, accessed on February 7, 2026,

https://pmc.ncbi.nlm.nih.gov/articles/PMC6604013/
7. Driver Distraction Using Visual-Based Sensors and Algorithms - PMC, accessed on February 7, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC5134464/
8. Identifying cognitive distraction using steering wheel reversal rates - White Rose Research Online, accessed on February 7, 2026, https://eprints.whiterose.ac.uk/id/eprint/102869/1/Reversal%20rates%20identify%20the%20type%20of%20driver%20distraction%20Resubmission_final.pdf
9. What is Horizontal Gaze Nystagmus Testing in a DUI Case?, accessed on February 7, 2026, https://www.duilexington.com/blog/2023/january/what-is-horizontal-gaze-nystagmus-testing-in-a-d/
10. Gaze entropy measures detect alcohol-induced driver impairment - ResearchGate, accessed on February 7, 2026, https://www.researchgate.net/publication/335387491_Gaze_entropy_measures_detect_alcohol-induced_driver_impairment
11. US8899748B1 - Automated detection of eye nystagmus - Google Patents, accessed on February 7, 2026, https://patents.google.com/patent/US8899748B1/en
12. Horizontal Gaze Nystagmus - The National College for DUI Defense, accessed on February 7, 2026, https://www.ncdd.com/dui-defenses-that-work/horizontal-gaze-nystagmus
13. Enhancing Road Safety with AI: A Secure System for Detecting Impairment through Horizontal Gaze Nystagmus - James Madison University, accessed on February 7, 2026, https://www.jmu.edu/economicdevelopment/_files/enhancing_road_safety_with_ai__a_secure__personalized_system_for_detecting_impairment_through_horizontal_gaze_nystagmus.pdf
14. A Pilot Study of Smartphone Eye Tracking for Detection of Positional Nystagmus - PubMed, accessed on February 7, 2026, https://pubmed.ncbi.nlm.nih.gov/40688101/
15. Field Sobriety Test Demonstration: What is Horizontal Gaze Nystagmus? - YouTube, accessed on February 7, 2026, https://www.youtube.com/watch?v=MZLh8YSCeSw
16. Field Sobriety Test Review | Texas District & County Attorneys Association, accessed on February 7, 2026, https://www.tdcaa.com/resources/dwi/field-sobriety-test-review/
17. Eye Gaze Entropy Reflects Individual Experience in the Context of Driving - MDPI, accessed on February 7, 2026, https://www.mdpi.com/1099-4300/28/1/8
18. Gaze Transition Entropy as a Measure of Attention Allocation in a Dynamic Workspace Involving Automation - ODU Digital Commons, accessed on February 7, 2026, https://digitalcommons.odu.edu/cgi/viewcontent.cgi?article=1213&context=psychology_fac_pubs
19. Remote Heart Rate Prediction in Virtual Reality Head-Mounted Displays Using Machine Learning Techniques - MDPI, accessed on February 7, 2026,

https://www.mdpi.com/1424-8220/22/23/9486
20. Contactless Camera-Based Heart Rate and Respiratory Rate Monitoring Using AI on Hardware - PMC, accessed on February 7, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10181491/
21. pavisj/rppg-pos - GitHub, accessed on February 7, 2026, https://github.com/pavisj/rppg-pos
22. Design, Implementation and Evaluation of a Real-Time Remote Photoplethysmography (rPPG) Acquisition System for Non-Invasive Vital Sign Monitoring - arXiv, accessed on February 7, 2026, https://arxiv.org/html/2508.18787v1
23. A comprehensive review of heart rate measurement using remote photoplethysmography and deep learning - PMC, accessed on February 7, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12181896/
24. Drunk Driver Detection Using Multiple Non-Invasive Biosignals - MDPI, accessed on February 7, 2026, https://www.mdpi.com/1424-8220/25/5/1281
25. A Computer Algorithm Can See When You're Drunk - Popular Science, accessed on February 7, 2026, https://www.popsci.com/computer-algorithm-can-tell-when-youre-drunk/
26. Auto industry deadlines loom for impaired-driver detection tech, U-M offers a low-cost solution - Electrical and Computer Engineering, accessed on February 7, 2026, https://ece.engin.umich.edu/stories/auto-industry-deadlines-loom-for-impaired-driver-detection-tech-u-m-offers-a-low-cost-solution
27. Online Detection of Driver Fatigue Using Steering Wheel Angles for Real Driving Conditions, accessed on February 7, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC5375781/
28. Feasibility Study on Detecting Drunk Driving by Driving Maneuver and Vehicle Behavior - Research, accessed on February 7, 2026, https://www-nrd.nhtsa.dot.gov/pdf/ESV/Proceedings/25/25ESV-000296.pdf
29. (PDF) A Steering Wheel Reversal Rate Metric for Assessing Effects of Visual and Cognitive Secondary Task Load - ResearchGate, accessed on February 7, 2026, https://www.researchgate.net/publication/325843375_A_Steering_Wheel_Reversal_Rate_Metric_for_Assessing_Effects_of_Visual_and_Cognitive_Secondary_Task_Load
30. THE EFFECTS OF FATIGUE AND ALCOHOL ON HIGHWAY SAFETY - ROSA P, accessed on February 7, 2026, https://rosap.ntl.bts.gov/view/dot/1285/dot_1285_DS1.pdf
31. Effects of Different Alcohol Dosages on Steering Behavior in Curve Driving - ResearchGate, accessed on February 7, 2026, https://www.researchgate.net/publication/326890039_Effects_of_Different_Alcohol_Dosages_on_Steering_Behavior_in_Curve_Driving
32. Validating lane drifts as a predictive measure of drug or sleepiness induced driving impairment - PMC, accessed on February 7, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC7036056/
33. Effects of acute alcohol consumption on measures of simulated driving: A

systematic review and meta-analysis - PubMed, accessed on February 7, 2026, https://pubmed.ncbi.nlm.nih.gov/28343124/
34. Effect of Alcohol Use on Accelerating and Braking Behaviors of Drivers - PubMed, accessed on February 7, 2026, https://pubmed.ncbi.nlm.nih.gov/31039040/
35. Multisensor Data Fusion in IoT Environments in Dempster–Shafer Theory Setting: An Improved Evidence Distance-Based Approach - MDPI, accessed on February 7, 2026, https://www.mdpi.com/1424-8220/23/11/5141
36. Dempster Shafer Sensor Fusion for Autonomously Driving Vehicles - Diva-portal.org, accessed on February 7, 2026, http://www.diva-portal.org/smash/get/diva2:931627/FULLTEXT01.pdf
37. Eye-Blink Parameters Detect On-Road Track-Driving Impairment Following Severe Sleep Deprivation - PMC - NIH, accessed on February 7, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC6760410/
38. Semantically-Enhanced Feature Extraction with CLIP and Transformer Networks for Driver Fatigue Detection - MDPI, accessed on February 7, 2026, https://www.mdpi.com/1424-8220/24/24/7948
39. ALCOHOL IMPAIRMENT Of PERFORMANCE ON STEERING AND DISCRETE TASKS IN A DRIVING SIMULATOR - ROSA P, accessed on February 7, 2026, https://rosap.ntl.bts.gov/view/dot/1095/dot_1095_DS1.pdf
40. Build a Self-Driving Car in Carla Simulator with Python (Step-by-Step) - Towards AI, accessed on February 7, 2026, https://towardsai.net/p/machine-learning/build-a-self-driving-car-in-carla-simulator-with-python-step-by-step
41. Development of a driving simulator for virtual experience and education on drunk driving, accessed on February 7, 2026, https://www.atsinternationaljournal.com/2011-issues/development-of-a-driving-simulator-for-virtual-experience-and-education-on-drunk-driving/
42. Car Simulation of Drunk Driving Behavior - LACCEI.org, accessed on February 7, 2026, https://laccei.org/LACCEI2025-Mexico/full-papers/Contribution_627_final_a.pdf
43. A Simulator Dataset to Support the Study of Impaired Driving - arXiv, accessed on February 7, 2026, https://arxiv.org/html/2507.02867v1
44. DIF: Dataset of perceived Intoxicated Faces for drunk person identification, accessed on February 7, 2026, https://research.monash.edu/en/publications/dif-dataset-of-perceived-intoxicated-faces-for-drunk-person-ident/
45. Advancing Intoxication Detection: A Smartwatch-Based Approach - arXiv, accessed on February 7, 2026, https://arxiv.org/html/2510.09916v1
46. Detecting Driver Drowsiness with Multi-Sensor Data Fusion Combined with Machine Learning - ROSA P, accessed on February 7, 2026, https://rosap.ntl.bts.gov/view/dot/59148/dot_59148_DS1.pdf
47. SparsePPG: Towards Driver Monitoring Using Camera-Based Vital Signs Estimation in Near-Infrared - IEEE Xplore, accessed on February 7, 2026, https://ieeexplore.ieee.org/document/8575330/
48. DUI (Driver Under Influence) Detection Using Open CV, accessed on February 7,

2026, https://irjaeh.com/index.php/journal/article/download/723/659/1441
49. ECU researchers develop visual tracking technology to detect drink drivers | EurekAlert!, accessed on February 7, 2026, https://www.eurekalert.org/news-releases/1049166
50. Carla Driver Behaviour Dataset - Kaggle, accessed on February 7, 2026, https://www.kaggle.com/datasets/dasmehdixtr/carla-driver-behaviour-dataset
51. Monitoring the effect of alcohol intake via facial temperature variations using thermography synchronized with heartbeat - PubMed Central, accessed on February 7, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12402095/
52. Development of non-invasive alcohol analyzer using Photoplethysmographytle, accessed on February 7, 2026, https://www.researchgate.net/publication/357657530_Development_of_non-invasive_alcohol_analyzer_using_Photoplethysmographytle
53. Next-Day Effects of Social Drinking on Driver Fatigue and Driving Performance - Diva-Portal.org, accessed on February 7, 2026, http://www.diva-portal.org/smash/get/diva2:1999139/FULLTEXT01.pdf
54. Prolonged nocturnal driving can be as dangerous as severe alcohol-impaired driving - PubMed, accessed on February 7, 2026, https://pubmed.ncbi.nlm.nih.gov/21226780/
55. Multi-Modal Fusion Transformer for End-to-End Autonomous Driving - YouTube, accessed on February 7, 2026, https://www.youtube.com/watch?v=WxadQyQ2gMs
56. Interaction-Based Drunk Driving Detection in a Real Vehicle Using Driver Monitoring Cameras and Rea - Research Collection, accessed on February 7, 2026, https://www.research-collection.ethz.ch/bitstreams/ff9a0429-6394-4d69-bb45-1a6b95c4f8b0/download
57. Python implementation of the Approximate Entropy cryptographic test for randomness, accessed on February 7, 2026, https://gist.github.com/StuartGordonReid/ff86c5a895fa90b0880e
58. Identifying Method of Drunk Driving Based on Driving Behavior - ResearchGate, accessed on February 7, 2026, https://www.researchgate.net/publication/233352798_Identifying_Method_of_Drunk_Driving_Based_on_Driving_Behavior
59. Drunk Driver Detection Using Thermal Facial Images - MDPI, accessed on February 7, 2026, https://www.mdpi.com/2078-2489/16/5/413